

▼ P3: GAN (30%)

In this problem, we will train a generative adversarial network (GAN) to generate new celebrities after showing it pictures of many real celebrities.

```

1 %matplotlib inline
2
3 from __future__ import print_function
4
5 import argparse
6 import os
7 import random
8 import torch
9 import torch.nn as nn
10 import torch.nn.parallel
11 import torch.backends.cudnn as cudnn
12 import torch.optim as optim
13 import torch.utils.data
14 import torchvision.datasets as dset
15 import torchvision.transforms as transforms
16 import torchvision.utils as vutils
17 import numpy as np
18 import matplotlib.pyplot as plt
19 import matplotlib.animation as animation
20 from IPython.display import HTML
21
22 # Set random seed for reproducibility
23 manualSeed = 999
24 # manualSeed = random.randint(1, 10000) # use if you want new results
25 print("Random Seed: ", manualSeed)
26 random.seed(manualSeed)
27 torch.manual_seed(manualSeed)

Random Seed:  999
<torch._C.Generator at 0x7dd8140fd250>

```

▼ (a) Prepare CelebA Dataset

we will use the [Celeb-A Faces dataset](#) which can be downloaded from the [link](#). The dataset will download as a file named `img_align_celeba.zip`. Once downloaded, create a directory named `CelebA` and extract the zip file into that directory. Then, set the `dataroot` input for this notebook to the `CelebA` directory you just created. The resulting directory structure should be:

```

./path/to/CelebA
    -> img_align_celeba
        -> 188242.jpg
        -> 173822.jpg
        -> 284702.jpg
        -> 537394.jpg
        ...

```

This is an important step because we will be using the `ImageFolder` dataset class, which requires there to be subdirectories in the dataset's root folder.

```

1 !gdown --fuzzy https://drive.google.com/file/d/0B7EVK8r0v71pZjFTYXZW3F1RnM/view?usp=share_link&resourcekey=0-dYn9z10tMJ0BAkviAcfdyQ img_align_celeba.zip
/bin/bash: line 1: img_align_celeba.zip: command not found
Downloading...
From: https://drive.google.com/uc?id=0B7EVK8r0v71pZjFTYXZW3F1RnM
To: /content/img_align_celeba.zip
100% 1.44G/1.44G [00:20<00:00, 71.0MB/s]

1 !mkdir ./CelebA
2 !unzip /content/img_align_celeba.zip -d ./CelebA

```

```
Set parameters for the implemented network.

1 # Root directory for dataset
2 dataroot = "./CelebA"
3 # Number of workers for dataloader
4 workers = 2
5 # Batch size during training
6 batch_size = 256
7 # Spatial size of training images. All images will be resized to this size using a transformer.
8 image_size = 64
9 # Number of channels in the training images. For color images this is 3
10 nc = 3
11 # Size of z latent vector (i.e. size of generator input)
12 nz = 100
13 # Learning rate for optimizers
14 lr = 0.0005
15 # Beta1 hyperparam for Adam optimizers
16 beta1 = 0.5
```

Now, we can create the dataset, create the dataloader, set the device to run on, and finally visualize some of the training data.

```

1 # We can use an image folder dataset the way we have it setup.
2 # Create the dataset
3 dataset = dset.ImageFolder(root=dataroot,
4                             transform=transforms.Compose([
5                                 transforms.Resize(image_size),
6                                 transforms.CenterCrop(image_size),
7                                 transforms.ToTensor(),
8                                 transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5)),
9                             ]))
10 # Create the dataloader
11 dataloader = torch.utils.data.DataLoader(dataset, batch_size=batch_size,
12                                         shuffle=True, num_workers=workers)
13
14 # Decide which device we want to run on
15 device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
16
17 # Plot some training images
18 real_batch = next(iter(dataloader))
19 plt.figure(figsize=(8,8))
20 plt.axis("off")
21 plt.title("Training Images")
22 plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[192:], padding=2, normalize=True).cpu(),(1,2,0)))

```

<matplotlib.image.AxesImage at 0x7dd72dc71a20>

Training Images



▼ (b) GAN Implementation

▼ 1. Weight Initialization

All model weights in GAN should be randomly initialized from a Normal distribution with `mean = 0`, `stdev = 0.02`. The `weights_init` function takes an initialized model as input and reinitializes all convolutional, convolutional-transpose, and batch normalization layers to meet this criteria. This function is applied to the models immediately after initialization.

```

1 # custom weights initialization called on netG and netD
2 def weights_init(m):
3     classname = m.__class__.__name__
4     if classname.find('Conv') != -1:
5         nn.init.normal_(m.weight.data, 0.0, 0.02)
6     elif classname.find('BatchNorm') != -1:
7         nn.init.normal_(m.weight.data, 1.0, 0.02)
8         nn.init.constant_(m.bias.data, 0)

```

▼ 2. Generator

The generator G , is designed to map the latent space vector z to data-space. Since our data are images, converting z to data-space means ultimately creating an RGB image with the same size as the training images (i.e. $3 \times 64 \times 64$). In practice, this is accomplished through a series of strided two dimensional convolutional transpose layers, each paired with a 2D batch norm layer and a relu activation. The output of the generator is fed through a tanh function to return it to the input data range of $[-1, 1]$.

```

1 class Generator(nn.Module):
2     def __init__(self):
3         super(Generator, self).__init__()
4         # TODO: finish implementing the network architecture
5         self.seq = nn.Sequential(
6             # TODO: the first ConvTranspose2d layer will take the noise as input with size of 100
7             # so you will need at least one ConvTranspose2d layer, one BatchNorm2d layer followed by a ReLU activation here
8             nn.ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False),
9             nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True),
10            nn.ReLU(),
11            # ENDS HERE
12            nn.ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False),
13            nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True),
14            nn.ReLU(),
15            nn.ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False),
16            nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True),
17            nn.ReLU(),
18            nn.ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False),
19            nn.BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True),
20            nn.ReLU(),
21            nn.ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False),
22            nn.Tanh()
23        )
24
25     def forward(self, input):
26         x = self.seq(input)
27         return x
28
29 # Create the generator
30 netG = Generator().to(device)
31
32 # Apply the weights_init function to randomly initialize all weights
33 # to mean=0, stddev=0.2.
34 netG.apply(weights_init)
35
36 # Print the model
37 print(netG)
38
39 Generator(
40     (seq): Sequential(
41         (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)
42         (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
43         (2): ReLU()
44         (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
45         (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
46         (5): ReLU()
47         (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
48         (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
49         (8): ReLU()
50         (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
51         (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
52         (11): ReLU()
53         (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
54         (13): Tanh()
55     )
56 )

```

```

Generator(
  (seq): Sequential(
    (0): ConvTranspose2d(100, 512, kernel_size=(4, 4), stride=(1, 1), bias=False)
    (1): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (2): ReLU()
    (3): ConvTranspose2d(512, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (4): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (5): ReLU()
    (6): ConvTranspose2d(256, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (7): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (8): ReLU()
    (9): ConvTranspose2d(128, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (10): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (11): ReLU()
    (12): ConvTranspose2d(64, 3, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (13): Tanh()
  )
)

```

3. Discriminator

The discriminator D is a binary classification network that takes an image as input and outputs a scalar probability that the input image is real (as opposed to fake). Here, D takes a $3 \times 64 \times 64$ input image, processes it through a series of Conv2d, BatchNorm2d, and ReLU layers, and outputs the final probability through a Sigmoid activation function. This architecture can be extended with more layers if necessary for the problem, but there is significance to the use of the strided convolution, BatchNorm, and ReLUs.

```

1 class Discriminator(nn.Module):
2     def __init__(self):
3         super(Discriminator, self).__init__()
4         # TODO: finish implementing the network architecture
5         self.seq = nn.Sequential(
6             nn.Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False),
7             nn.ReLU(),
8             nn.Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False),
9             nn.BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True),
10            nn.ReLU(),
11            nn.Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False),
12            nn.BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True),
13            nn.ReLU(),
14            nn.Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False),
15            nn.BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True),
16            nn.ReLU(),
17            # TODO: the last Conv2d layer of the doscriminator followed by a sigmoid
18            nn.Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), padding=(0, 0), bias=False),
19            nn.Sigmoid()
20            # ENDS HERE
21        )
22
23     def forward(self, input):
24         x = self.seq(input)
25         return x
26
27 # Create the Discriminator
28 netD = Discriminator().to(device)
29
30 # Apply the weights_init function to randomly initialize all weights
31 # to mean=0, stdev=0.2.
32 netD.apply(weights_init)
33
34 # Print the model
35 print(netD)
36
Discriminator(
  (seq): Sequential(
    (0): Conv2d(3, 64, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (1): ReLU()
    (2): Conv2d(64, 128, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (3): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (4): ReLU()
    (5): Conv2d(128, 256, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (6): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (7): ReLU()
    (8): Conv2d(256, 512, kernel_size=(4, 4), stride=(2, 2), padding=(1, 1), bias=False)
    (9): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
    (10): ReLU()
    (11): Conv2d(512, 1, kernel_size=(4, 4), stride=(1, 1), bias=False)
  )
)

```

```
(12): Sigmoid()
)
)
```

▼ Loss Functions and Optimizers

With D and G setup, we can specify how they learn through the loss functions and optimizers. We will use the Binary Cross Entropy loss ([BCELoss](#)) function which is defined in PyTorch as:

$$\ell(x, y) = L = \{l_1, \dots, l_n\}^\top, \quad l_n = -[y_n \cdot \log x_n + (1 - y_n) \cdot \log(1 - x_n)]$$

Notice how this function provides the calculation of both log components in the objective function (i.e. $\log(D(x))$ and $\log(1 - D(G(z)))$). We can specify what part of the BCE equation to use with the y input. This is accomplished in the training loop which is coming up soon, but it is important to understand how we can choose which component we wish to calculate just by changing y (i.e. GT labels).

```
1 # Initialize BCELoss function
2 criterion = nn.BCELoss()
3
4 # Create batch of latent vectors that we will use to visualize
5 # the progression of the generator
6 fixed_noise = torch.randn(64, nz, 1, 1, device=device)
7
8 # Establish convention for real and fake labels during training
9 real_label = 1.
10 fake_label = 0.
11
12 # Setup Adam optimizers for both G and D
13 optimizerD = optim.Adam(netD.parameters(), lr=lr, betas=(beta1, 0.999))
14 optimizerG = optim.Adam(netG.parameters(), lr=lr, betas=(beta1, 0.999))
```

▼ 5. Training Loop

Finally, now that we have all of the parts of the GAN framework defined, we can train it. Be mindful that training GANs is somewhat of an art form, as incorrect hyperparameter settings lead to mode collapse with little explanation of what went wrong. Here, we will closely follow Algorithm 1 from Goodfellow's paper, while abiding by some of the best practices shown in [gan_hacks](#). Namely, we will construct different mini-batches for real and fake images, and also adjust G 's objective function to maximize $\log(D(G(z)))$. Training is split up into two main parts. Part 1 updates the Discriminator and Part 2 updates the Generator.

Part 1: Train the Discriminator

Recall, the goal of training the discriminator is to maximize the probability of correctly classifying a given input as real or fake. In terms of Goodfellow, we wish to "update the discriminator by ascending its stochastic gradient". Practically, we want to maximize $\log(D(x)) + \log(1 - D(G(z)))$. Due to the separate mini-batch suggestion from gan_hacks, we will calculate this in two steps. First, we will construct a batch of real samples from the training set, forward pass through D , calculate the loss ($\log(D(x))$), then calculate the gradients in a backward pass. Secondly, we will construct a batch of fake samples with the current generator, forward pass this batch through D , calculate the loss ($\log(1 - D(G(z)))$), and accumulate the gradients with a backward pass. Now, with the gradients accumulated from both the all-real and all-fake batches, we call a step of the Discriminator's optimizer.

Part 2: Train the Generator

As stated in the original paper, we want to train the Generator by minimizing $\log(1 - D(G(z)))$ in an effort to generate better fakes. As mentioned, this was shown by Goodfellow to not provide sufficient gradients, especially early in the learning process. As a fix, we instead wish to maximize $\log(D(G(z)))$. In the code we accomplish this by: classifying the Generator output from Part 1 with the Discriminator, computing G 's loss using *real labels* as *GT*, computing G 's gradients in a backward pass, and finally updating G 's parameters with an optimizer step. It may seem counter-intuitive to use the real labels as *GT* labels for the loss function, but this allows us to use the $\log(x)$ part of the BCELoss (rather than the $\log(1 - x)$ part) which is exactly what we want.

Finally, we will do some statistic reporting and at the end of each epoch we will push our `fixed_noise` batch through the generator to visually track the progress of G 's training. The training statistics reported are:

- **Loss_D** - discriminator loss calculated as the sum of losses for the all real and all fake batches ($\log(D(x)) + \log(1 - D(G(z)))$).
- **Loss_G** - generator loss calculated as $\log(D(G(z)))$
- **D(x)** - the average output (across the batch) of the discriminator for the all real batch. This should start close to 1 then theoretically converge to 0.5 when G gets better. Think about why this is.
- **D(G(z))** - average discriminator outputs for the all fake batch. The first number is before D is updated and the second number is after D is updated. These numbers should start near 0 and converge to 0.5 as G gets better. Think about why this is.

Note: This step might take a while, depending on how many epochs you run and if you removed some data from the dataset.

```

1 # Training Loop
2 num_epochs = 30
3
4 # Lists to keep track of progress
5 img_list = []
6 G_losses = []
7 D_losses = []
8 iters = 0
9
10 print("Starting Training Loop...")
11 # For each epoch
12 for epoch in range(num_epochs):
13     # For each batch in the dataloader
14     for i, data in enumerate(dataloader, 0):
15         #####
16         # 1. Update D network: maximize log(D(x)) + log(1 - D(G(z)))
17         #####
18
19         # Initialize the gradient of D
20         netD.zero_grad()
21
22         # 1.1 Calculate the loss with all-real batch
23
24         # Format batch
25         real_batch = data[0].to(device)
26         b_size = real_batch.size(0)
27         label = torch.full((b_size,), real_label, dtype=torch.float, device=device)
28         # Forward pass real batch through D
29         output = netD(real_batch).view(-1)
30         # Calculate loss on all-real batch
31         errD_real = criterion(output, label)
32         # Calculate gradients for D in backward pass
33         errD_real.backward()
34         D_x = output.mean().item()
35
36         # 1.2 Calculate the loss with all-fake batch
37         # Generate batch of latent vectors
38         noise = torch.randn(b_size, nz, 1, 1, device=device)
39         # Generate fake image batch with G and a label array
40         fake = netG(noise)
41         label.fill_(fake_label)
42         # Classify all fake batch with D
43         output = netD(fake.detach()).view(-1)
44         # Calculate D's loss on the all-fake batch
45         errD_fake = criterion(output, label)
46         # Calculate the gradients for this batch
47         errD_fake.backward()
48         # Calculate the mean of output, D_G_z1, from netD with your fake batch
49         D_G_z1 = output.mean().item()
50
51         # 1.3 Combine the loss from all-real batch and all-fake batch
52         # Add the gradients from the all-real and all-fake batches
53         errD = errD_real + errD_fake
54         # Update D with step()
55         optimizerD.step()
56
57         #####
58         # 2. Update G network: maximize log(D(G(z)))
59         #####
60
61         # TODO: Initialize the gradient of G
62         netG.zero_grad()
63         # TODO: create the label array, remember fake labels are real for generator cost
64
65         b_size = fake.size(0)
66         label = torch.full((b_size,), real_label, dtype=torch.float, device=device)
67
68         # TODO: Since we just updated D, perform another forward pass of all-fake batch through D as output (you are not suppose to detach)
69         output = netD(fake).view(-1)
70
71         # TODO: Calculate G's loss based on this output
72         errG = criterion(output, label)
73         # TODO: Calculate gradients using backward() for G's loss
74         errG.backward()
75         # TODO: Calculate the mean of output, D_G_z2, from netD with your fake batch

```

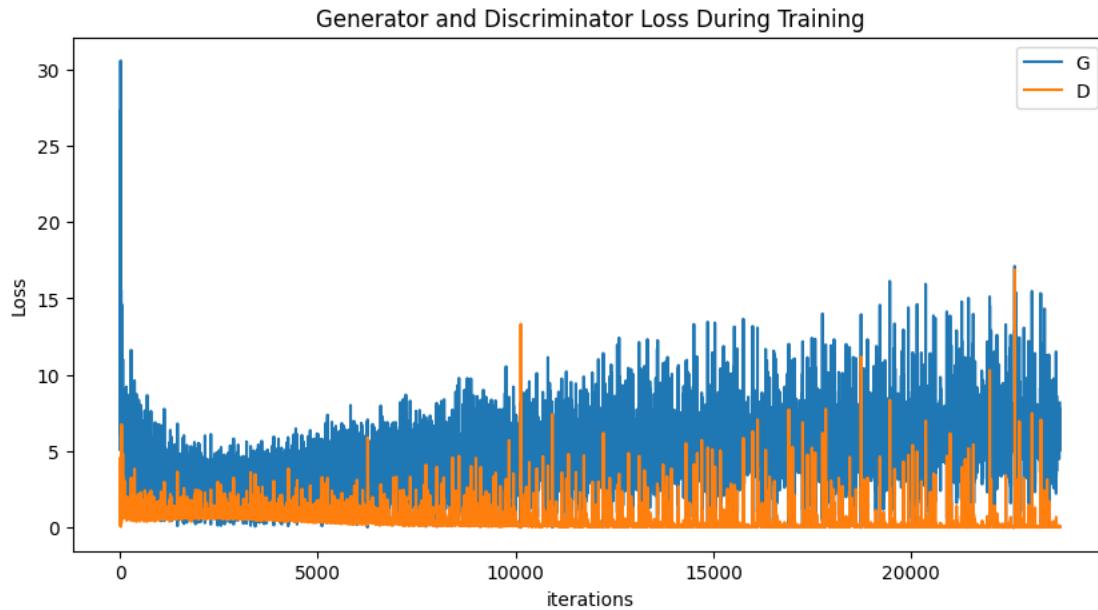
```

76     D_G_z2 = output.mean().item()
77     # TODO: Update G with step
78     optimizerG.step()
79     # Output training stats
80     if i % 50 == 0:
81         print('[%d/%d][%d/%d]\tLoss_D: %.4f\tLoss_G: %.4f\tD(x): %.4f\tD(G(z)): %.4f / %.4f'
82             % (epoch, num_epochs, i, len(dataloader),
83                 errD.item(), errG.item(), D_x, D_G_z1, D_G_z2))
84
85     # Save Losses for plotting later
86     G_losses.append(errG.item())
87     D_losses.append(errD.item())
88
89     # Check how the generator is doing by saving G's output on fixed_noise
90     if (iters % 500 == 0) or ((epoch == num_epochs-1) and (i == len(dataloader)-1)):
91         with torch.no_grad():
92             fake = netG(fixed_noise).detach().cpu()
93             img_list.append(vutils.make_grid(fake, padding=2, normalize=True))
94
95     iters += 1

```

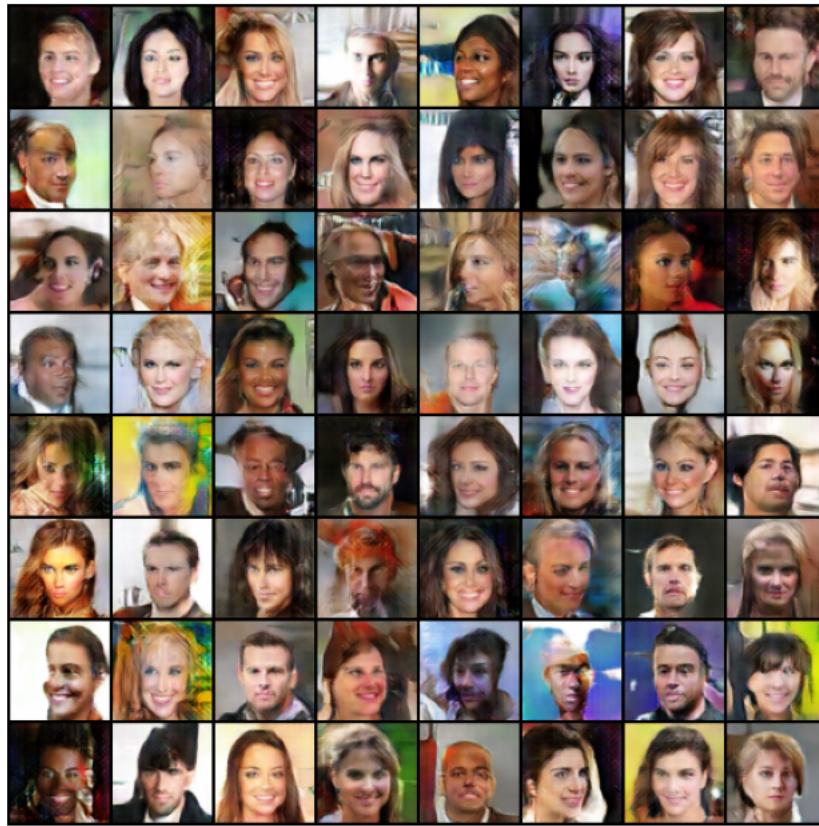
[23/30][550/792] Loss_D: 0.6572 Loss_G: 2.1628 D(x): 0.7085 D(G(z)): 0.1237 / 0.2095
[23/30][600/792] Loss_D: 0.1144 Loss_G: 5.1625 D(x): 0.9280 D(G(z)): 0.0272 / 0.0153
[23/30][650/792] Loss_D: 0.0487 Loss_G: 7.0937 D(x): 0.9931 D(G(z)): 0.0356 / 0.0035
[23/30][700/792] Loss_D: 0.0517 Loss_G: 5.3933 D(x): 0.9803 D(G(z)): 0.0271 / 0.0169
[23/30][750/792] Loss_D: 0.1566 Loss_G: 3.6292 D(x): 0.8818 D(G(z)): 0.0105 / 0.0869
[24/30][0/792] Loss_D: 0.1715 Loss_G: 7.8478 D(x): 0.9992 D(G(z)): 0.1277 / 0.0011
[24/30][50/792] Loss_D: 0.2565 Loss_G: 10.5653 D(x): 0.9965 D(G(z)): 0.1685 / 0.0001
[24/30][100/792] Loss_D: 0.1906 Loss_G: 5.2642 D(x): 0.8647 D(G(z)): 0.0140 / 0.0201
[24/30][150/792] Loss_D: 0.0520 Loss_G: 6.1463 D(x): 0.9602 D(G(z)): 0.0077 / 0.0090
[24/30][200/792] Loss_D: 0.0603 Loss_G: 7.4105 D(x): 0.9957 D(G(z)): 0.0451 / 0.0027
[24/30][250/792] Loss_D: 0.0478 Loss_G: 7.7187 D(x): 0.9716 D(G(z)): 0.0128 / 0.0021
[24/30][300/792] Loss_D: 0.0760 Loss_G: 5.9679 D(x): 0.9731 D(G(z)): 0.0381 / 0.0096
[24/30][350/792] Loss_D: 0.0390 Loss_G: 5.6607 D(x): 0.9772 D(G(z)): 0.0136 / 0.0141
[24/30][400/792] Loss_D: 0.0471 Loss_G: 7.3052 D(x): 0.9879 D(G(z)): 0.0217 / 0.0027
[24/30][450/792] Loss_D: 0.1189 Loss_G: 7.8220 D(x): 0.9831 D(G(z)): 0.0777 / 0.0013
[24/30][500/792] Loss_D: 0.0984 Loss_G: 4.7855 D(x): 0.9871 D(G(z)): 0.0656 / 0.0259
[24/30][550/792] Loss_D: 0.0776 Loss_G: 6.2984 D(x): 0.9850 D(G(z)): 0.0511 / 0.0064
[24/30][600/792] Loss_D: 0.0662 Loss_G: 5.3100 D(x): 0.9600 D(G(z)): 0.0171 / 0.0212
[24/30][650/792] Loss_D: 0.0700 Loss_G: 5.7274 D(x): 0.9800 D(G(z)): 0.0180 / 0.0123
[24/30][700/792] Loss_D: 0.0644 Loss_G: 5.7618 D(x): 0.9978 D(G(z)): 0.0516 / 0.0098
[24/30][750/792] Loss_D: 0.0333 Loss_G: 5.5257 D(x): 0.9794 D(G(z)): 0.0103 / 0.0147
[25/30][0/792] Loss_D: 0.2920 Loss_G: 9.9751 D(x): 0.9996 D(G(z)): 0.1975 / 0.0002
[25/30][50/792] Loss_D: 0.0255 Loss_G: 7.8301 D(x): 0.9818 D(G(z)): 0.0062 / 0.0018
[25/30][100/792] Loss_D: 0.0511 Loss_G: 5.4476 D(x): 0.9714 D(G(z)): 0.0183 / 0.0152
[25/30][150/792] Loss_D: 0.2562 Loss_G: 6.2525 D(x): 0.9790 D(G(z)): 0.1592 / 0.0061
[25/30][200/792] Loss_D: 0.0481 Loss_G: 6.0783 D(x): 0.9762 D(G(z)): 0.0194 / 0.0098
[25/30][250/792] Loss_D: 0.5461 Loss_G: 5.3578 D(x): 0.9328 D(G(z)): 0.2860 / 0.0125
[25/30][300/792] Loss_D: 0.0339 Loss_G: 7.2994 D(x): 0.9924 D(G(z)): 0.0233 / 0.0034
[25/30][350/792] Loss_D: 0.6465 Loss_G: 14.6206 D(x): 0.9993 D(G(z)): 0.3711 / 0.0000
[25/30][400/792] Loss_D: 0.1198 Loss_G: 5.0480 D(x): 0.9842 D(G(z)): 0.0828 / 0.0171
[25/30][450/792] Loss_D: 0.0429 Loss_G: 5.8961 D(x): 0.9881 D(G(z)): 0.0267 / 0.0110
[25/30][500/792] Loss_D: 0.0735 Loss_G: 5.8975 D(x): 0.9917 D(G(z)): 0.0576 / 0.0083
[25/30][550/792] Loss_D: 0.0660 Loss_G: 5.5604 D(x): 0.9585 D(G(z)): 0.0182 / 0.0139
[25/30][600/792] Loss_D: 0.5295 Loss_G: 7.4288 D(x): 0.9977 D(G(z)): 0.3192 / 0.0016
[25/30][650/792] Loss_D: 0.1196 Loss_G: 6.7669 D(x): 0.9941 D(G(z)): 0.0786 / 0.0036
[25/30][700/792] Loss_D: 0.0926 Loss_G: 4.7999 D(x): 0.9559 D(G(z)): 0.0337 / 0.0278
[25/30][750/792] Loss_D: 0.0481 Loss_G: 5.6837 D(x): 0.9892 D(G(z)): 0.0325 / 0.0128
[26/30][0/792] Loss_D: 0.0289 Loss_G: 5.8287 D(x): 0.9965 D(G(z)): 0.0241 / 0.0094
[26/30][50/792] Loss_D: 0.0313 Loss_G: 7.3741 D(x): 0.9762 D(G(z)): 0.0056 / 0.0032
[26/30][100/792] Loss_D: 0.0412 Loss_G: 5.9588 D(x): 0.9771 D(G(z)): 0.0145 / 0.0102
[26/30][150/792] Loss_D: 0.2954 Loss_G: 2.4196 D(x): 0.7947 D(G(z)): 0.0009 / 0.2300
[26/30][200/792] Loss_D: 0.0374 Loss_G: 5.8779 D(x): 0.9788 D(G(z)): 0.0139 / 0.0108
[26/30][250/792] Loss_D: 0.0668 Loss_G: 5.0814 D(x): 0.9584 D(G(z)): 0.0160 / 0.0234
[26/30][300/792] Loss_D: 0.0239 Loss_G: 6.1323 D(x): 0.9846 D(G(z)): 0.0073 / 0.0082
[26/30][350/792] Loss_D: 0.0388 Loss_G: 6.0153 D(x): 0.9913 D(G(z)): 0.0279 / 0.0075
[26/30][400/792] Loss_D: 0.3139 Loss_G: 3.6840 D(x): 0.9221 D(G(z)): 0.1485 / 0.0709
[26/30][450/792] Loss_D: 0.1083 Loss_G: 5.8080 D(x): 0.9776 D(G(z)): 0.0694 / 0.0083
[26/30][500/792] Loss_D: 0.0298 Loss_G: 6.8103 D(x): 0.9926 D(G(z)): 0.0204 / 0.0041
[26/30][550/792] Loss_D: 0.0507 Loss_G: 5.8690 D(x): 0.9919 D(G(z)): 0.0376 / 0.0089
[26/30][600/792] Loss_D: 0.0423 Loss_G: 6.2575 D(x): 0.9716 D(G(z)): 0.0109 / 0.0096
[26/30][650/792] Loss_D: 0.0148 Loss_G: 7.1708 D(x): 0.9930 D(G(z)): 0.0067 / 0.0049
[26/30][700/792] Loss_D: 0.0484 Loss_G: 6.3147 D(x): 0.9971 D(G(z)): 0.0406 / 0.0079
[26/30][750/792] Loss_D: 0.0553 Loss_G: 7.2559 D(x): 0.9545 D(G(z)): 0.0049 / 0.0039
[27/30][0/792] Loss_D: 0.4406 Loss_G: 9.6543 D(x): 0.9924 D(G(z)): 0.2608 / 0.0002
[27/30][50/792] Loss_D: 0.0737 Loss_G: 6.7291 D(x): 0.9981 D(G(z)): 0.0604 / 0.0042
[27/30][100/792] Loss_D: 0.1161 Loss_G: 7.3114 D(x): 0.9826 D(G(z)): 0.0698 / 0.0026
[27/30][150/792] Loss_D: 0.0695 Loss_G: 6.7972 D(x): 0.9976 D(G(z)): 0.0530 / 0.0040
[27/30][200/792] Loss_D: 0.2250 Loss_G: 5.6607 D(x): 0.9677 D(G(z)): 0.1399 / 0.0080
[27/30][250/792] Loss_D: 0.0449 Loss_G: 6.9443 D(x): 0.9892 D(G(z)): 0.0290 / 0.0043

```
1
2 plt.figure(figsize=(10,5))
3 plt.title("Generator and Discriminator Loss During Training")
4 plt.plot(G_losses,label="G")
5 plt.plot(D_losses,label="D")
6 plt.xlabel("iterations")
7 plt.ylabel("Loss")
8 plt.legend()
9 plt.show()
```



```
1 fig = plt.figure(figsize=(8,8))
2 plt.axis("off")
3 ims = [[plt.imshow(np.transpose(i,(1,2,0)), animated=True)] for i in img_list]
4 ani = animation.ArtistAnimation(fig, ims, interval=1000, repeat_delay=1000, blit=True)
5
6 HTML(ani.to_jshtml())
```

WARNING:matplotlib.animation:Animation size has reached 21210989 bytes, exceeding the limit of 20971520.0. If you're sure you want a la



```

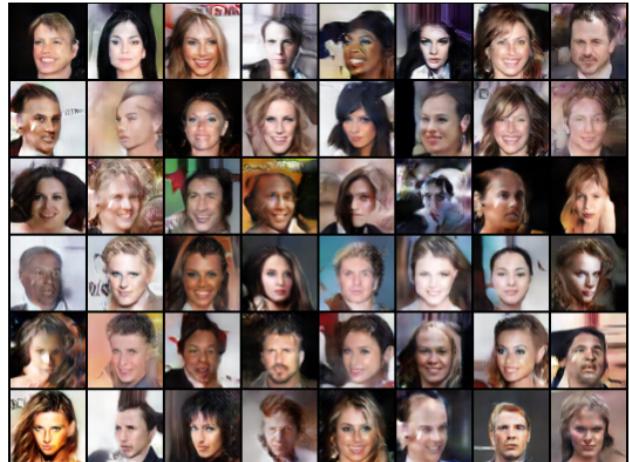
1 # Grab a batch of real images from the dataloader
2 real_batch = next(iter(dataloader))
3
4 # Plot the real images
5 plt.figure(figsize=(15,15))
6 plt.subplot(1,2,1)
7 plt.axis("off")
8 plt.title("Real Images")
9 plt.imshow(np.transpose(vutils.make_grid(real_batch[0].to(device)[:, :64], padding=5, normalize=True).cpu(), (1,2,0)))
10
11 # Plot the fake images from the last epoch
12 plt.subplot(1,2,2)
13 plt.axis("off")
14 plt.title("Fake Images")
15 plt.imshow(np.transpose(img_list[-1], (1,2,0)))
16 plt.show()

```

Real Images



Fake Images



1

